

# Neural Network-Based PDE Solver Using the Feynman-Kac Formula

Shubha Sanket Samantaray

*Université Grenoble Alpes*

November 27, 2025

# Motivation

---

- Partial Differential Equations (PDEs) are fundamental tools for modeling dynamic systems in physics, biology, finance, and engineering.
- Classical solvers such as finite differences and finite elements are effective in low dimensions but become computationally intractable in high-dimensional settings due to the curse of dimensionality.
- Find an alternate method to solve PDEs in a mesh-free setting, being able to perform well in higher dimensions as well.
- Feynman–Kac formula reformulates certain linear parabolic PDEs as expectations over stochastic processes, from deterministic solvers to statistical estimation.

# The Heat Equation: A Foundational Example

---

- Backward form (terminal condition):

$$\partial_t u + \frac{1}{2} \Delta u = 0, \quad u(T, x) = f(x)$$

- Feynman–Kac solution:

$$u(t, x) = \mathbb{E}[f(W_T) \mid W_t = x]$$

- Forward form (initial condition):

$$\partial_t v - \frac{1}{2} \Delta v = 0, \quad v(0, x) = f(x)$$

- Feynman–Kac solution:

$$v(t, x) = \mathbb{E}^x[f(W_t)] = \mathbb{E}[f(W_t) \mid W_0 = x]$$

- $W_t$  denotes a standard Brownian motion in  $\mathbb{R}^d$ .

# Feynman–Kac for Linear Parabolic PDEs

---

- Consider the PDE:

$$\partial_t u + b(x)\partial_x u + \frac{1}{2}\sigma^2(x)\partial_{xx}u - ru = 0, \quad u(T, x) = f(x)$$

- **Feynman–Kac representation:**

$$u(t, x) = \mathbb{E} \left[ e^{-r(T-t)} f(X_T) \mid X_t = x \right]$$

- **Where**  $X_s$  satisfies the SDE:

$$dX_s = b(X_s) ds + \sigma(X_s) dW_s$$

- Connects PDE solutions to stochastic processes through conditional expectations.

# The Black–Scholes Model

---

- The Black–Scholes model provides a framework for pricing European options.
- A European call option gives its holder, the right (but not the obligation) to buy an asset at strike price  $K$  at maturity  $T$ .
- The option price  $V(t, S)$  satisfies the celebrated Black–Scholes PDE:

$$\partial_t V + rS \partial_S V + \frac{1}{2} \sigma^2 S^2 \partial_{SS} V - rV = 0,$$

with terminal condition:

$$V(T, S) = \Phi(S) = \max(S - K, 0).$$

- This PDE models the evolution of option prices under a no-arbitrage assumption and constant volatility.

# Feynman–Kac Representation of Black–Scholes

---

- The Black–Scholes PDE is a special case of a linear parabolic PDE:

$$\partial_t u + b(x)\partial_x u + \frac{1}{2}\sigma^2(x)\partial_{xx} u - ru = 0$$

with:

$$x = S, \quad b(S) = rS, \quad \sigma(S) = \sigma S.$$

- Under the risk-neutral measure  $\mathbb{Q}$ , the asset follows:

$$dS_t = rS_t dt + \sigma S_t dW_t^{\mathbb{Q}}.$$

- Feynman–Kac gives the option price as:

$$V(t, S) = \mathbb{E}^{\mathbb{Q}} \left[ e^{-r(T-t)} \Phi(S_T) \mid S_t = S \right].$$

- In particular, for a European call:

$$C(t, S) = \mathbb{E}^{\mathbb{Q}} \left[ e^{-r(T-t)} (S_T - K)^+ \mid S_t = S \right].$$

# Neural Networks as Conditional Expectation Approximators

---

- The Feynman–Kac formula expresses the PDE solution as a conditional expectation:

$$u(t, x) = \mathbb{E} \left[ e^{-r(T-t)} f(X_T) \mid X_t = x \right]$$

- This is the solution to the regression problem, with the MSE loss:

$$\varphi^* = \arg \min_{\varphi} \mathbb{E} \left[ (g(X_T) - \varphi(X_t))^2 \right]$$


- The unique minimizer is the conditional expectation:

$$\varphi^*(x) = \mathbb{E} [g(X_T) \mid X_t = x]$$

- Note that in this setup, we have to train the model by keeping the time  $t$  fixed.

# PDEs, Probabilities, and Neural Networks: A Conceptual Map

A conceptual flow: from mathematical theory to data-driven approximation.



# Learning Procedure

---

1. **Sample Generation:** Simulate trajectories of the underlying SDE to obtain pairs  $(X_t^i, X_T^i)$ .
2. **Label Construction:** Evaluate the terminal payoff to compute training targets:  $y^i = f(X_T^i)$ .
3. **Model Training:** Fit a neural network  $\varphi_\theta$  by minimizing the mean squared error between  $\varphi_\theta(X_t^i)$  and  $y^i$ .
4. **Solution Recovery:** Use the trained model  $\varphi_\theta(x) \approx \mathbb{E}[f(X_T) | X_t = x]$  as an approximation of the PDE solution.

# Data Generation for European Call Option

---

Two Sampling Strategies:

- Closed-form (GBM):

$$S_T = S_t \exp \left( \left( r - \frac{1}{2} \sigma^2 \right) (T - t) + \sigma \sqrt{T - t} Z \right), \quad Z \sim \mathcal{N}(0, 1)$$

- Euler–Maruyama discretization:

$$S_{n+1} = S_n + rS_n \Delta t + \sigma S_n \Delta W_n, \quad \Delta W_n \sim \mathcal{N}(0, \Delta t)$$

**Note:** This dual-scheme setup lets us assess accuracy and robustness under both exact and approximate SDE sampling.

# Results for the Forward Heat Equation

---




Figure: Left: Scatter plot of Brownian pairs  $(W_0, W_1)$ . Right: Histogram of target values  $W_1^2$ .

**Observation:** The histogram is positively skewed, consistent with the square of a Gaussian variable.

# Heat Equation: Training and Prediction Behavior




Figure: Left: Training loss curve. Right: Predicted values vs stochastic labels  $W_1^2$ .

**Observation:** Despite label noise, the network learns the conditional mean, not just data interpolation.

# Heat Equation: Comparison with Exact Solution

- Overlay shows excellent alignment between predicted and exact solutions.
- The absolute pointwise error is uniformly small, staying below 0.05 throughout the domain.




Figure: Left: Predicted function vs exact  $x^2 + 1$ . Right: Absolute error across domain.

**Conclusion:** The Feynman–Kac solver accurately recovers the heat equation solution and validates the regression framework.

# Black–Scholes: Euler–Maruyama Sampling

- Asset paths simulated via Euler–Maruyama discretization of GBM.
- Network trained for 5000 epochs to learn pricing map at  $t = 0.5$ .
- Predicted prices show close alignment with the analytical Black–Scholes solution.





Figure: Left: Training loss convergence. Right: NN predictions vs analytical Black–Scholes prices.

# Black–Scholes: Closed-Form Sampling

- Terminal stock prices generated via exact GBM formula.
- Neural network again trained for 5000 epochs.
- Even better approximation observed compared to Euler–Maruyama.





Figure: Left: Training loss for closed-form data. Right: NN predictions vs Black–Scholes values.

# Black–Scholes: Discussion and Insights

---

- Both data generation schemes produce accurate neural approximations of the Black–Scholes price function at fixed time  $t = 0.5$ .
- **Closed-form simulation** yields lower approximation error:
  - Mean absolute error: 0.001085
  - Maximum error: 0.003068
- **Euler–Maruyama discretization** performs slightly worse:
  - Mean absolute error: 0.001482
  - Maximum error: 0.005502
- The comparison confirms that while both methods validate the neural solver, the exact simulation offers a quantitative edge when available.

**Conclusion:** Closed-form and numerical schemes both validate the neural Feynman–Kac solver. Euler–Maruyama remains a robust choice when exact simulations are unavailable.

# Learning the Full Space-Time Solution Surface

**Goal:** Extend the fixed-time model  $x \mapsto V(t, x)$  to a full surface model

$$(t, x) \mapsto V(t, x) = \mathbb{E}^{\mathbb{Q}} \left[ e^{-r(T-t)} \Phi(S_T) \mid S_t = x \right]$$

One network  $\varphi_{\theta}(t, x)$  learns the entire solution surface across  $t \in [0, T]$  and  $x \in \mathbb{R}_+$ .

## Motivation:

- In fixed- $t$  experiments, networks learn a slice  $x \mapsto V(t, x)$ .
- This requires retraining for each new  $t$ .
- We aim to extend the input to include both  $t$  and  $x$ , enabling generalization.

**Key Question:** Can a neural network learn a map of conditional expectations over time and space, i.e. a continuous family of such expectations?

# Training Setup for Space-Time Approximation

---

## Data Generation Procedure:

- Sample  $(t^i, x^i) \in [0, T] \times [a, b]$  uniformly or on a grid.
- Simulate terminal value:

$$S_T^i = x^i \cdot \exp \left( \left( r - \frac{1}{2}\sigma^2 \right) (T - t^i) + \sigma \sqrt{T - t^i} \cdot Z^i \right), \quad Z^i \sim \mathcal{N}(0, 1)$$

- Compute label:

$$y^i = e^{-r(T-t^i)} \Phi(S_T^i)$$

## Learning Objective:

$$\min_{\theta} \frac{1}{N} \sum_{i=1}^N (\varphi_{\theta}(t^i, x^i) - y^i)^2$$

# Learning the Full Surface





Figure: Left: Neural network surface  $V_{\text{NN}}(t, x)$ . Right: Exact Black–Scholes solution  $V_{\text{BS}}(t, x)$ .

# Training and Accuracy Evaluation

- Loss stabilizes during training, confirming convergence.
- Learned surface matches analytical solution with low pointwise error.




Figure: Left: Training loss over epochs. Right: Absolute error across the full  $(t, x)$  domain.

**Conclusion:** The 2D network successfully generalizes over both time and space.

# Comparison with 1D Network and Classical Solvers

- Slice extracted at  $t = 0.5$  from the full surface.
- Compares 1D vs 2D networks — both learn the solution accurately.
- All solvers (NNs, MC, FD) show strong agreement with the exact solution.





Figure: Left: 1D NN vs 2D NN vs exact at  $t = 0.5$ . Right: All methods vs exact at  $t = 0.5$ .

# Comparison with 1D Network and Classical Solvers

- Left: 2D network is compared against Monte Carlo simulation at  $t = 0.5$ .
- Right: Comparison with Crank–Nicolson finite difference scheme.
- Neural network predictions show excellent alignment with reference methods.





Figure: Left: 1D NN, 2D NN vs Monte Carlo. Right: 1D NN, 2D NN vs Finite Difference.

# Comparison of Evaluation Methods

---

- **Monte Carlo:** Computes expectation at each  $(t, x)$  via fresh simulations. No reuse across inputs.
- **1D Neural Network:** Learns  $x \mapsto V(t, x)$  for fixed  $t$ . Requires retraining for each new time.
- **Finite Difference:** Solves on a fixed time-space grid. Interpolation needed for off-grid points.
- **2D Neural Network:** Learns the full map  $(t, x) \mapsto V(t, x)$ . Predicts values instantly for any input pair.

## Key Insight

**2D neural networks enable mesh-free, reusable inference across all  $(t, x)$  values**  
— *a significant advantage over simulation and grid-based methods.*

# Conclusion

---

- Reformulated linear parabolic PDEs as conditional expectation problems using the Feynman–Kac formula.
- Trained neural networks on simulated SDE trajectories to approximate the PDE solutions.
- Successfully recovered known solutions to the heat equation and Black–Scholes PDE at fixed time and across full surfaces.
- Achieved strong agreement with finite difference and Monte Carlo methods.
- Enabled mesh-free, data-driven PDE solvers that scale to irregular or high-dimensional settings.

# Future Work

---

- **High-Dimensional Extensions:** Apply the solver to high-dimensional PDEs where classical methods fail.
- **Nonlinear PDEs:** Extend the method using backward SDEs and stochastic control techniques.
- **Variance Reduction:** Incorporate techniques like importance sampling to improve training efficiency.
- **Financial Applications:** Test the method on exotic and path-dependent derivatives for real-world relevance.

# Thank You

---

Questions?